

Dog Breeds Classification

This notebook contains the code for training a dog breeds classification model on the dog-breed-image-dataset using PyTorch.

- This is a small-scale model trained on only 967 images and contains only 10 classes.
- link for the dataset: <https://www.kaggle.com/datasets/khushikhushikhushi/dog-breed-image-dataset>

0. Device agnostic code

```
In [1]: import torch
        from torchvision import datasets, transforms
        import os
```

```
In [2]: device= 'cuda' if torch.cuda.is_available() else 'cpu'
        device
```

```
Out[2]: 'cuda'
```

1. loading the data

```
In [3]: data_path='/kaggle/input/dog-breed-image-dataset/dataset'
```

```
In [4]: transform = transforms.Compose([
        transforms.Resize((128, 128)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])
```

```
In [5]: dataset=datasets.ImageFolder(data_path, transform=transform)
```

```
In [6]: classes=dataset.classes
        print(f'There are {len(classes)} classes in the dataset set as:\n{classes}')
```

There are 10 classes in the dataset set as:

['Beagle', 'Boxer', 'Bulldog', 'Dachshund', 'German_Shepherd', 'Golden_Retrieve
r', 'Labrador_Retriever', 'Poodle', 'Rottweiler', 'Yorkshire_Terrier']

```
In [7]: data=[]
        for img, lbl in dataset:
            data.append((img, lbl))
        print(f'The length of the dataset: {len(data)}')
```

The length of the dataset: 967

1.2 Visualizing the images

```
In [8]: import matplotlib.pyplot as plt
import numpy as np
import random


index = random.randint(0, len(data) - 1)

img, lbl = data[index]

image = img.numpy().transpose((1, 2, 0))

plt.imshow(image)
plt.title(f'Label: {classes[lbl]}\n Shape: {image.shape}')
plt.axis('off')
plt.show()
```

Label: Boxer
Shape: (128, 128, 3)



2 loading the data into DataLoader

```
In [9]: from sklearn.model_selection import train_test_split
tensors, labels=zip(*data)
tensor=np.array([t.numpy() for t in tensors])
label=np.array([l for l in labels])
X_train, X_valid, y_train, y_valid = train_test_split(tensor, label, test_size=0
```

```
In [10]: import torch
from torch.utils.data import TensorDataset, DataLoader
```

```

X_train_tensor = torch.tensor(X_train, dtype=torch.float32)
y_train_tensor = torch.tensor(y_train, dtype=torch.long)
X_valid_tensor = torch.tensor(X_valid, dtype=torch.float32)
y_valid_tensor = torch.tensor(y_valid, dtype=torch.long)

train_dataset = TensorDataset(X_train_tensor, y_train_tensor)
valid_dataset = TensorDataset(X_valid_tensor, y_valid_tensor)

train_loader = DataLoader(dataset=train_dataset, batch_size=32, shuffle=True, dr
valid_loader = DataLoader(dataset=valid_dataset, batch_size=32, shuffle=False, d

for batch in train_loader:
    inputs, labels = batch
    print(inputs.shape, labels.shape)
    break

```

```
torch.Size([32, 3, 128, 128]) torch.Size([32])
```

3 Building a model

```

In [11]: import torch
import torch.nn as nn
import torch.nn.functional as F

class model_1(nn.Module):
    def __init__(self):
        super(model_1, self).__init__()

        self.block_1= nn.Sequential(
            nn.Conv2d(in_channels=3, out_channels=16, kernel_size=3,
            nn.ReLU(),
            nn.Conv2d(in_channels=16, out_channels=32, kernel_size=3
            nn.ReLU(),
            nn.Conv2d(in_channels=32, out_channels=64, kernel_size=3
            nn.ReLU(),
            nn.MaxPool2d(kernel_size=2, stride=2, padding=0)
        )
        self.block_2= nn.Sequential(
            nn.Flatten(),
            nn.Linear(in_features=64 * 64 * 64, out_features=128),
            nn.Dropout(p=0.5),
            nn.Linear(in_features=128, out_features=10)
        )
    def forward(self, x):
        x=self.block_1(x)
        x=self.block_2(x)
        return x

model = model_1().to(device)

print(model)

```

```
model_1(
  (block_1): Sequential(
    (0): Conv2d(3, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (1): ReLU()
    (2): Conv2d(16, 32, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (3): ReLU()
    (4): Conv2d(32, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))
    (5): ReLU()
    (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  )
  (block_2): Sequential(
    (0): Flatten(start_dim=1, end_dim=-1)
    (1): Linear(in_features=262144, out_features=128, bias=True)
    (2): Dropout(p=0.5, inplace=False)
    (3): Linear(in_features=128, out_features=10, bias=True)
  )
)
```

```
In [12]: try:
          import torchinfo
        except:
          !pip install torchinfo
          import torchinfo

          from torchinfo import summary
          print(f'The summary of the model.')
          summary(model, input_size=[1, 3, 128, 128]) # do a test pass through of an example
```

The summary of the model.

```

Out[12]: =====
=====
Layer (type:depth-idx)                Output Shape                Param #
=====
=====
model_1                                [1, 10]                     --
├─Sequential: 1-1                      [1, 64, 64, 64]            --
│   └─Conv2d: 2-1                      [1, 16, 128, 128]          448
│       └─ReLU: 2-2                    [1, 16, 128, 128]          --
│           └─Conv2d: 2-3                [1, 32, 128, 128]          4,640
│               └─ReLU: 2-4              [1, 32, 128, 128]          --
│                   └─Conv2d: 2-5          [1, 64, 128, 128]          18,496
│                       └─ReLU: 2-6        [1, 64, 128, 128]          --
│                           └─MaxPool2d: 2-7 [1, 64, 64, 64]            --
├─Sequential: 1-2                      [1, 10]                     --
│   └─Flatten: 2-8                     [1, 262144]                --
│       └─Linear: 2-9                   [1, 128]                    33,554,560
│           └─Dropout: 2-10              [1, 128]                    --
│               └─Linear: 2-11           [1, 10]                     1,290
=====
=====
Total params: 33,579,434
Trainable params: 33,579,434
Non-trainable params: 0
Total mult-adds (M): 419.96
=====
=====
Input size (MB): 0.20
Forward/backward pass size (MB): 14.68
Params size (MB): 134.32
Estimated Total Size (MB): 149.20
=====
=====

```

4. defining functions like earlystopping, train_step, test_step and train.

```

In [13]: class Earlystopping:
          def __init__(self, filepath, patience=10, minimize=True):
              self.filepath=filepath
              self.patience=patience
              self.minimize=minimize
              self.counter=0
              self.best_score=None
              self.early_stop=False

              if not os.path.exists(filepath):
                  os.makedirs(filepath)
                  print(f'file path did not exist so {filepath} created')
              else:
                  print(f'Directory exists {filepath}')

          def __call__(self, val_metric, model):

              if self.minimize:
                  score=-val_metric
              else:

```

```

        score=val_metric

        if self.best_score is None:
            self.best_score=score
        elif score< self.best_score:
            self.counter+=1
            if self.counter>= self.patience:
                print(f'Early Stopping')
                self.early_stop=True
            else:
                self.best_score=score
                self.counter=0

    def save_checkpoint(self, model):
        print(f'Saving model with the current best validation metric: {self.best
        torch.save(model.state_dict(), os.path.join(self.filepath, 'checkpoint.p

```

```

In [14]: def train_step(model,dataloader, loss_fn,optimizer):
    model.train()
    train_loss, train_acc = 0, 0
    for (X,y) in dataloader:
        X,y=X.to(device),y.to(device)

        y_pred=model(X)
        loss=loss_fn(y_pred, y)
        train_loss+=loss.item()
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        y_pred_class = torch.argmax(torch.softmax(y_pred, dim=1), dim=1)
        train_acc +=(y_pred_class == y).sum().item()/len(y_pred)

    train_loss = train_loss / len(dataloader)
    train_acc = train_acc / len(dataloader)
    return train_loss, train_acc

```

```

In [15]: def valid_step(model, dataloader, loss_fn):
    model.eval() # Set the model to evaluation mode
    total_loss, total_acc = 0, 0
    num_batches = len(dataloader)

    with torch.no_grad(): # Disable gradient calculations
        for X, y in dataloader:
            X, y = X.to(device), y.to(device)

            # Forward pass
            y_pred = model(X)

            # Loss calculation
            loss = loss_fn(y_pred, y) # Use logits (y_pred) and class indices (
            total_loss += loss.item()

            # Accuracy calculation
            y_pred_class = torch.argmax(torch.softmax(y_pred, dim=1), dim=1)
            total_acc +=(y_pred_class == y).sum().item()/len(y_pred) # Compute

    # Average loss and accuracy over all batches

```

```

avg_loss = total_loss / num_batches
avg_acc = total_acc / num_batches

return avg_loss, avg_acc

```

```

In [16]: from tqdm.auto import tqdm

early_stopping=Earlystopping(filepath='/kaggle/working/checkpoints', patience=1)
def train(epochs,model,train_dataloader,valid_dataloader,loss_fn,optimizer):
    results={'training_loss':[],'training_acc':[],'validation_loss':[],'validation_acc':[]}

    for epoch in tqdm(range(epochs)):
        train_loss,train_acc=train_step(model,train_dataloader,loss_fn,optimizer)
        valid_loss,valid_acc=valid_step(model,valid_dataloader,loss_fn)

        print(
            f"Epoch: {epoch+1} | "
            f"train_loss: {train_loss:.4f} | "
            f"train_acc: {train_acc:.4f} | "
            f"valid_loss: {valid_loss:.4f} | "
            f"valid_acc: {valid_acc:.4f}"
        )
        results["training_loss"].append(train_loss)
        results["training_acc"].append(train_acc)
        results["validation_loss"].append(valid_loss)
        results["validation_acc"].append(valid_acc)
        early_stopping(valid_loss, model)
        if early_stopping.early_stop:
            early_stopping.save_checkpoint(model)
            print(f'Early stopping triggered')
            break

    return results

```

Directory exists /kaggle/working/checkpoints

5. Model training

```

In [17]: torch.manual_seed(42)
torch.cuda.manual_seed(42)
NUM_EPOCHS = 50
loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(params=model.parameters(), lr=0.01)
from timeit import default_timer as timer
start_time = timer()

model_results = train(model=model,
                      train_dataloader=train_loader,
                      valid_dataloader=valid_loader,
                      optimizer=optimizer,
                      loss_fn=loss_fn,
                      epochs=NUM_EPOCHS)

end_time = timer()
print(f"Total training time: {end_time-start_time:.3f} seconds")

```

0%| | 0/50 [00:00<?, ?it/s]

Epoch: 1		train_loss: 2.1987		train_acc: 0.2122		valid_loss: 2.0421		valid_acc: 0.3698
Epoch: 2		train_loss: 1.7287		train_acc: 0.4036		valid_loss: 1.6869		valid_acc: 0.4271
Epoch: 3		train_loss: 1.3010		train_acc: 0.5898		valid_loss: 0.9707		valid_acc: 0.7292
Epoch: 4		train_loss: 0.8826		train_acc: 0.7096		valid_loss: 0.7190		valid_acc: 0.8021
Epoch: 5		train_loss: 0.6684		train_acc: 0.8060		valid_loss: 0.5975		valid_acc: 0.8229
Epoch: 6		train_loss: 0.5602		train_acc: 0.8372		valid_loss: 0.4117		valid_acc: 0.9219
Epoch: 7		train_loss: 0.2152		train_acc: 0.9466		valid_loss: 0.5103		valid_acc: 0.8594
Epoch: 8		train_loss: 0.1794		train_acc: 0.9622		valid_loss: 0.1152		valid_acc: 0.9792
Epoch: 9		train_loss: 0.1089		train_acc: 0.9714		valid_loss: 0.0750		valid_acc: 0.9896
Epoch: 10		train_loss: 0.1381		train_acc: 0.9596		valid_loss: 0.1246		valid_acc: 0.9792
Epoch: 11		train_loss: 0.1012		train_acc: 0.9688		valid_loss: 0.0915		valid_acc: 0.9896
Epoch: 12		train_loss: 0.0429		train_acc: 0.9883		valid_loss: 0.0785		valid_acc: 0.9896
Epoch: 13		train_loss: 0.0730		train_acc: 0.9805		valid_loss: 0.2570		valid_acc: 0.9219
Epoch: 14		train_loss: 0.0590		train_acc: 0.9883		valid_loss: 0.0563		valid_acc: 0.9896
Epoch: 15		train_loss: 0.0453		train_acc: 0.9883		valid_loss: 0.0673		valid_acc: 0.9896
Epoch: 16		train_loss: 0.0156		train_acc: 0.9961		valid_loss: 0.0599		valid_acc: 0.9896
Epoch: 17		train_loss: 0.0121		train_acc: 0.9987		valid_loss: 0.1397		valid_acc: 0.9792
Epoch: 18		train_loss: 0.0095		train_acc: 0.9987		valid_loss: 0.0676		valid_acc: 0.9896
Epoch: 19		train_loss: 0.0044		train_acc: 1.0000		valid_loss: 0.0680		valid_acc: 0.9896
Epoch: 20		train_loss: 0.0116		train_acc: 0.9961		valid_loss: 0.0516		valid_acc: 0.9896
Epoch: 21		train_loss: 0.0049		train_acc: 0.9987		valid_loss: 0.0565		valid_acc: 0.9896
Epoch: 22		train_loss: 0.0085		train_acc: 0.9987		valid_loss: 0.0500		valid_acc: 0.9896
Epoch: 23		train_loss: 0.0038		train_acc: 1.0000		valid_loss: 0.0698		valid_acc: 0.9896
Epoch: 24		train_loss: 0.0035		train_acc: 1.0000		valid_loss: 0.0677		valid_acc: 0.9896
Epoch: 25		train_loss: 0.0186		train_acc: 0.9974		valid_loss: 0.0400		valid_acc: 0.9896
Epoch: 26		train_loss: 0.0050		train_acc: 1.0000		valid_loss: 0.0604		valid_acc: 0.9896
Epoch: 27		train_loss: 0.0057		train_acc: 0.9987		valid_loss: 0.0597		valid_acc: 0.9896
Epoch: 28		train_loss: 0.0024		train_acc: 1.0000		valid_loss: 0.0521		valid_acc: 0.9948
Epoch: 29		train_loss: 0.0026		train_acc: 1.0000		valid_loss: 0.0511		valid_acc: 0.9896
Epoch: 30		train_loss: 0.0013		train_acc: 1.0000		valid_loss: 0.0527		valid_acc: 0.9896

```
Epoch: 31 | train_loss: 0.0012 | train_acc: 1.0000 | valid_loss: 0.0570 | valid_acc: 0.9896
Epoch: 32 | train_loss: 0.0012 | train_acc: 1.0000 | valid_loss: 0.0619 | valid_acc: 0.9896
Epoch: 33 | train_loss: 0.0012 | train_acc: 1.0000 | valid_loss: 0.0607 | valid_acc: 0.9896
Epoch: 34 | train_loss: 0.0015 | train_acc: 1.0000 | valid_loss: 0.0579 | valid_acc: 0.9896
Epoch: 35 | train_loss: 0.0010 | train_acc: 1.0000 | valid_loss: 0.0591 | valid_acc: 0.9896
Early Stopping
Saving model with the current best validation metric: -0.04000071714593408
Early stopping triggered
Total training time: 23.573 seconds
```

6. Model testing on google images

```
In [18]: checkpoint_path = '/kaggle/working/checkpoints/checkpoint.pt'
model.load_state_dict(torch.load(checkpoint_path, map_location=device)) # Load

device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
model.to(device)
device
```

```
Out[18]: device(type='cuda')
```

```
In [19]: import requests
from PIL import Image
from io import BytesIO
from torchvision import transforms

def image_transform(image_url):

    response = requests.get(image_url)
    img = Image.open(BytesIO(response.content))

    transform = transforms.Compose([
        transforms.Resize((128, 128)),
        transforms.ToTensor(),
        transforms.Normalize(mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.225])
    ])

    input_tensor = transform(img)
    input_tensor = input_tensor.unsqueeze(0)

    return input_tensor
```

Please insert here the link for the image

```
In [20]: image_url='https://t4.ftcdn.net/jpg/00/86/67/45/360_F_86674540_3e6FWLLSXXKkjC5gw'
```

```
In [21]: image=image_transform(image_url)
```

```
In [22]: image_np=torch.squeeze(image, dim=0)
image_np = image_np.numpy().transpose((1, 2, 0))
```

```
plt.imshow(image_np )
plt.title(f'The url has this image init\n Shape: {image_np .shape}')
plt.axis('off')
plt.show()
```

The url has this image init
Shape: (128, 128, 3)



```
In [23]: sample = torch.unsqueeze(image, dim=0).to(device)
sample = torch.squeeze(sample, dim=1)
sample.shape
```

```
Out[23]: torch.Size([1, 3, 128, 128])
```

```
In [24]: with torch.no_grad():
    predictions = model(sample)
    pred_prob = torch.softmax(predictions.squeeze(), dim=0)
    pred_classes = pred_prob.argmax(dim=0)

    print(f'The image provided contains a {classes[pred_classes]}.')
```

The image provided contains a Bulldog.